

# Unsupervised Learning: clustering

AI534

## **Key Concepts**

Distance measures

Hierarchical Clustering: complete, single & average link algorithm

K-means algorithm

Mixture of Gaussians



Expectation maximization

Model selection of # of clusters

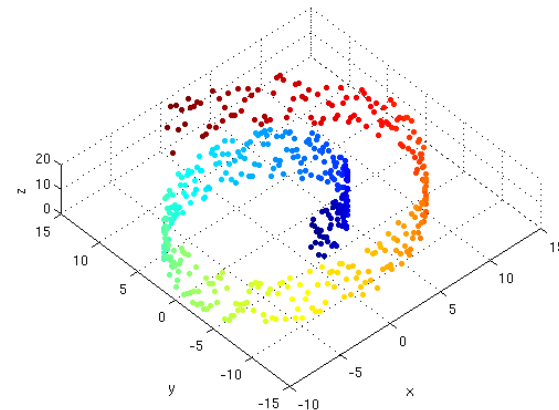
Evaluation of clustering

# What can we learn from unlabeled data?

- Finding Group of clusters in the data
- Finding low dimensional representation – dimension reduction

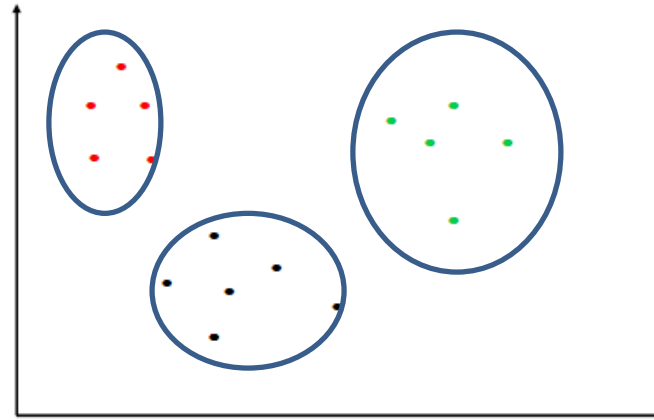


PCA



Nonlinear embedding

# Clustering



- Clustering: the process of grouping a set of objects into classes of similar objects
  - high within-class similarity
  - low cross-class similarity
- It is the most common form of unsupervised learning

# Example Applications

- Find genes that are similar in their functions
- Group documents based on topics
- Categorize customers based on their purchase history
- Group images based on their contents
- .....

# Critical Issues in Clustering

- What makes objects similar?
  - Definition of "similarity/distance"
- How many clusters?
  - Fixed a priori?
  - Completely data driven?
  - Avoid "trivial" clusters - too large or small
- Clustering Algorithms
  - Flat vs hierarchical
  - Hard vs soft

# What is similarity



Hard to define but  
We know it when we see it

- The real meaning of similarity is a philosophical question. We will take a more pragmatic approach
  - Depends on representation and algorithm. For many rep./alg., easier to think in terms of a distance (rather than similarity) between vectors

# What properties should a distance measure have?

- $D$  must be Symmetric
  - $D(A, B) = D(B, A)$
  - Otherwise, we can say  $A$  looks like  $B$  but  $B$  does not look like  $A$
- Positivity, and self-similarity
  - $D(A, B) \geq 0$ , and  $D(A, B) = 0$  iff  $A = B$
  - Otherwise there will be different objects that we cannot tell apart
- Must satisfy triangle inequality
  - $D(A, B) + D(B, C) \geq D(A, C)$
  - Otherwise one can say “ $A$  is like  $B$ ,  $B$  is like  $C$ , but  $A$  is not like  $C$  at all”

# Distance Measures: Minkowski Metric

- Suppose two object  $x$  and  $y$  both have  $d$  features
  - $x = (x_1, \dots, x_d), y = (y_1, \dots, y_d)$

- The Minkowski metric of order  $r$  is defined by

$$d(x, y) = \sqrt[r]{\sum_i |x_i - y_i|^r}$$

- Common Minkowski metrics:

- Euclidean( $r=2$ ):  $d(x, y) = \sqrt{\sum_i (x_i - y_i)^2}$ , also called  $L_2$  distance
- Manhattan distance( $r=1$ ) :  $d(x, y) = \sum_i |x_i - y_i|$ , also called  $L_1$  distance
- “Sup” distance( $r = +\infty$ ):  $d(x, y) = \max_i |x_i - y_i|$ , also called  $L_\infty$  distance



# Other Distances


- Hamming distance (Manhattan distance on binary features)
  - # of features that differ
  - e.g.: distance of two sites based on their species composition

	sp1	sp2	sp3	sp4	sp5	sp6	sp7	sp8	sp9
Site A:	1	0	1	1	0	0	1	0	1
Site B:	0	0	1	0	1	1	1	0	1

$$D(A, B) = 4$$

- Mahalanobis distance (assuming  $\mathbf{x}, \mathbf{y}$  follows a Gaussian distribution with covariance matrix  $\Sigma$ )

$$D(\mathbf{x}, \mathbf{y}) = \sqrt{(\mathbf{x} - \mathbf{y})^T \Sigma^{-1} (\mathbf{x} - \mathbf{y})}$$



# Similarities

- Cosine similarities – commonly used to measure document similarity


$$\cos(\mathbf{x}, \mathbf{x}') = \frac{\langle \mathbf{x} \cdot \mathbf{x}' \rangle}{|\mathbf{x}| \cdot |\mathbf{x}'|}$$

- Kernels – e.g., RBF (Gaussian) Kernel


$$S(X, X') = \exp \frac{-|X - X'|^2}{2\sigma^2}$$

# Clustering algorithms

- Hierarchical algorithms (not covered)
  - Bottom up – agglomerative
  - Top down – divisive

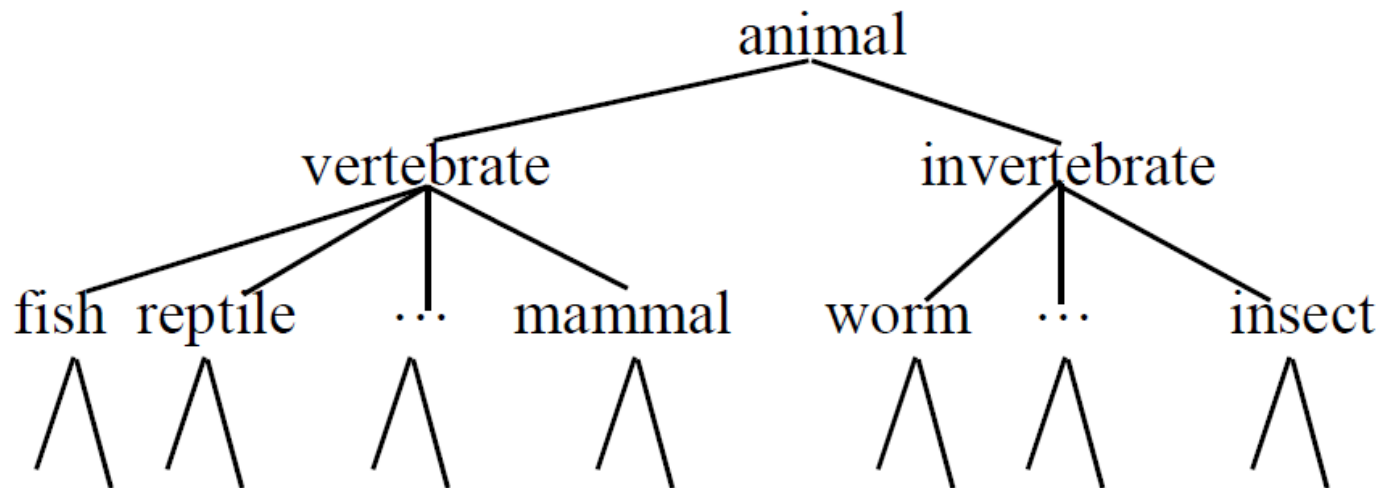


- Flat clustering algorithms
  - K-means
  - Mixture of Gaussian
  - Spectral Clustering (not covered)



# Hierarchical Clustering

- Given a set of objects, build a tree-based taxonomy



- Hierarchies are convenient way for organizing information

# Hierarchical Agglomerative Clustering (HAC)

- Starts with each object in a separate cluster
- While there are more than 1 cluster:
  - Find the **closest** pair of clusters and join them

The history of merging forms a tree of hierarchy




**Question**: how to measure the “**closeness**” of two clusters?

- Different definition leads to different algorithms



# Distance between clusters

The distance between two clusters is defined as follows


- Single-link
  - Distance between the nearest pair of points across two clusters:  $D(C_i, C_j) = \min_{x \in C_i, y \in C_j} d(x, y)$
- Complete-link
  - Distance between the furthest pair of points across two clusters:  $D(C_i, C_j) = \max_{x \in C_i, y \in C_j} d(x, y)$
- Average-link
  - Average distance of all cross-cluster pairs
  - $D(C_i, C_j) = \text{average}_{x \in C_i, y \in C_j} d(x, y)$
- Centroid
  - Distance between the means of the two clusters
  - $D(C_i, C_j) = d(\hat{x}, \hat{y})$




# Single Link




(1)





(2)




(3)




# Complete link




(1)



(2)




(3)



# Visualization: Dendrogram


- Height of the joint = the distance between the two merge clusters
- The merge distance monotonically increases as we merge more and more for
  - Single, complete and average linkage methods
  - But not for the centroid method



This example is shown upside down.




# Interpreting Dendrogram







- Dendrogram can be used to visually identify
  - the number of clusters in data
    - A horizontal cut creates a partition
    - Moving the cut from root down creates more clusters
    - Large gaps indicate good cutting points
  - well-formed clusters
    - Some clusters are better formed than the other
    - This can be easy to see on a dendrogram

# Example





# Single Link vs. Complete Link

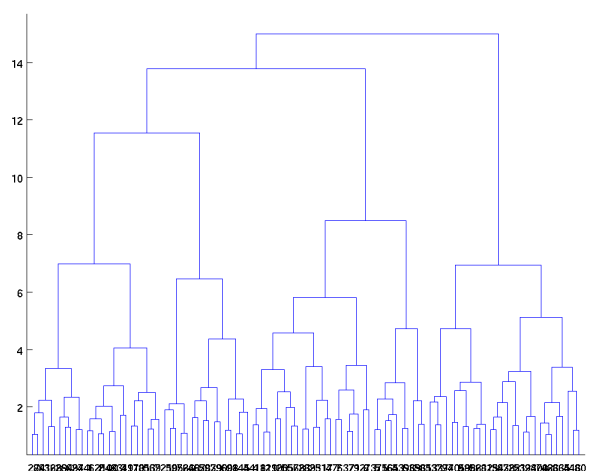



Which is single link? Which is complete link?

# Single Link vs. Complete Link



**Single**



**Complete**

- Single-link creates straggly clusters due to chaining effect

# Flat Clustering

- Given  $D$ , a data set of  $n$  points, we want to find  $k$  clusters
- For each cluster  $i$  we would like to represent all points of that cluster with one representative  $\mu_i$ , and the total representation error should be minimized:

$$\min_{\mu, C} \sum_{i=1}^k \sum_{x \in C_i} |x - \mu_i|^2$$

where  $C = \{C_1, C_2, \dots, C_k\}$  defines a partition of  $D$  such that

$$C_i \cap C_j = \emptyset, \forall i \neq j \text{ and } D = \cup_i C_i$$

and  $\mu_i$  is the representative of  $C_i$

- A combinatorial optimization problem
  - Discrete solution space
  - Exhaustive search for an optimal solution is not feasible

# An Iterative Solution

- ***Initialization:*** Start with a random partition of the data, or a random initial  $\mu'_i$ s
- ***Iterative step:*** update the cluster assignments and cluster centers to improve the objective
- ***Stopping criterion:*** if no improvement can be achieved.

# K-Means

## Algorithm

**Input** –  $N$  data points, and desired # of clusters:  $k$

**Initialize** –  $\mu_1, \dots, \mu_k$ , the  $k$  cluster centers (by randomly selecting  $k$  points)

**Iterate** –


1. Assigning each of the  $N$  data points to the closest  $\mu_i$
2. Re-estimate the cluster center by assuming that the current assignment is correct

$$\mu_i = \frac{1}{|C_i|} \sum_{\mathbf{x} \in C_i} \mathbf{x}$$

**Termination** –

If none of the data points changed membership in the last iteration, exit.  
Otherwise, go to 1

# Demo: K-Means Example (K=2)



Pick seeds

Reassign clusters

Compute centroids

Reassign clusters

Compute centroids

Reassign clusters

**Converged!**



# Does KMeans Converge?

## Objective

$$\min_{\mu, C} \sum_{i=1}^k \sum_{\mathbf{x} \in C_i} |\mathbf{x} - \mu_i|^2$$

1. Fix  $\mu_i$ 's, optimize  $C$ , the assignments

- Assign each point to its closet center

***Step 1 of kmeans***

2. Fix  $C$ , optimize  $\mu$ :

$$\min_{\mu} \sum_{i=1}^k \sum_{\mathbf{x} \in C_i} |\mathbf{x} - \mu_i|^2$$

- Take partial derivative of  $\mu_i$  and set to zero, we have


$$2 \sum_{\mathbf{x} \in C_i} (\mathbf{x} - \mu_i) = 0 \Rightarrow \mu_i = \frac{1}{|C_i|} \sum_{\mathbf{x} \in C_i} \mathbf{x}$$

***Step 2 of kmeans***

Each iteration is guaranteed to decrease the objective and there are finite possible  $C$ 's – thus guaranteed to converge --- but only to local minimum

# Impact of Initial Seeds

- Highly sensitive to the initial seeds



- Multiple random trials: choose the one with best sum of squared loss (important!)
- Heuristics for choosing better centers
  - choose initial centers to be far apart – furthest first traversal (kmeans ++)

# More Comments

- K-Means is exhaustive:
  - Cluster every data point, no notion of outlier
  - Outliers cause problems
    - Become singular clusters
    - Bias the centroid estimation
- K-medoids methods is more robust to outliers
  - Cluster medoid: must be one of the data points, one that has the minimum sum squared distance to all data points in the cluster
  - More expensive to compute
    - For each pt: sum squared dist with all other pts in cluster  $O(|C|^2)$
- Need to specify  $k$ : difficult in practice, many heuristic approaches

# Hard vs. Soft Clustering

- Hard clustering:
  - Data point is deterministically assigned to one and only one cluster
  - But in reality clusters may overlap
- Soft-clustering:
  - Data points are assigned to clusters with certain probabilities
- Often referred to as model-based clustering for the use of probabilistic model for clusters
  - Assuming each cluster follows a certain probabilistic (e.g., Gaussian) distribution

# Side track: Gaussian Bayes Classifier

- We have  $k$  classes in our data
- Each class contains data generated from a particular Gaussian distribution  $N(\mu_c, \Sigma_c), c = 1, \dots, k$
- The data is generated as follows:
  - randomly draw  $y \sim p(y)$
  - Randomly draw  $\mathbf{x} \sim N(\mu_y, \Sigma_y)$ , i.e., from the Gaussian distribution of class  $y$
  - Repeat  $n$  times to generate  $n$  points  $(\mathbf{x}_i, y_i)$
- In supervised learning, we observe  $(\mathbf{x}_i, y_i)$  in pairs and we simply need to estimate  $P(y = c)$  and  $p(\mathbf{x}|y = c)$  for  $c = 1, \dots, k$

# MLE Estimates of Mean and Covariance for Gaussian

Given that  $p(\mathbf{x}|y = c) \sim N(\mu_c, \Sigma_c)$ , and given a set of  $n_c$  training examples with  $y = c$ :

$$\mu_c = \frac{1}{n_c} \sum_{y_i=c} \mathbf{x}_i$$

$$\Sigma_c = \frac{1}{n_c} \sum_{y_i=c} (\mathbf{x}_i - \mu_c) (\mathbf{x}_i - \mu_c)^T$$

# Back to Unsupervised Learning

- Now assume we know our data is generated in the same way
- But for unsupervised learning, we don't have the  $y$ 's, and only observe  $\mathbf{x}$ 's
  - We observe a mixture of Gaussians (or mixture of other distributions)
- How can we learn the correct model from the incomplete data?
- We will still use Maximum likelihood estimation

# Gaussian Mixture Model

$$\begin{aligned} P(\mathbf{x}) &= \sum_{i=1}^k P(\mathbf{x}, y = i) \\ &= \sum_{i=1}^k P(\mathbf{x} | y = i) P(y = i) \\ &= \sum_{i=1}^k \alpha_i P(\mathbf{x} | \theta_i) \end{aligned}$$

$\alpha_i = P(y = i)$ : class priors  
Mixing parameter

$\theta_i = \{\mu_i, \Sigma_i\}$

Goal of learning:

- Given a set of  $\mathbf{x}$ 's, find the values of  $\{\alpha_1, \dots, \alpha_k, \theta_1, \dots, \theta_k\}$  that maximize the likelihood of observing the  $\mathbf{x}$ 's



# Maximum Marginal Likelihood

$$\begin{aligned}\arg \max_{\theta} \prod_j P(\mathbf{x}^j) &= \arg \max_{\theta} \prod_j \sum_{i=1}^k P(\mathbf{x}^j, y^j = i) \\ &= \arg \max_{\theta} \sum_{j=1}^n \underbrace{\log \sum_{i=1}^k P(\mathbf{x}^j, y^j = i)}\end{aligned}$$


log sum is difficult to optimize !

Gradient ascent is doable but very inefficient

# Expectation Maximization (EM)

- Commonly used approach for dealing with hidden (missing) data
  - Here the cluster labels are hidden
- Iterative algorithm that starts with an initial guess of the model parameters
- Iteratively performs two linked steps:
  - **Expectation (E-step)**: given current model parameters  $\lambda_t$ , compute the expectation for the hidden (missing) data
  - **Maximization (M-step)**: re-estimate the parameters  $\lambda_{t+1}$  assuming that the expected values computed in the E-step are the true values
- We will first show how it works for mixture of Gaussian

# EM – simple case



- A simple case: spherical Gaussians
  - We have unlabeled data  $x^1, \dots, x^m$
  - We know there are  $K$  classes
  - We know  $\alpha_1 = P(y = 1), \dots, \alpha_K = P(y = K)$
  - We don't know  $\mu_1 \dots \mu_K$ , but know the common variance  $\sigma^2$

Start with an initial guess for  $\mu_1, \dots, \mu_K$ ,

1. If we know  $\mu_1, \dots, \mu_K$ , we can easily compute probability that a point  $x^j$  belongs to class  $k$ :

$$P(y = k | x^j) \propto \exp\left(-\frac{1}{2\sigma^2} |x^j - \mu_k|^2\right) p(y = k)$$

Simply evaluate this, then normalize


**E-step**

2. If we know *the* probability that each point belongs to each class, we can estimate the  $\mu_1, \dots, \mu_K$

$$\mu_k = \frac{\sum_{j=1}^m P(y=k|x^j) x^j}{\sum_{j=1}^m p(y=k|x^j)}$$

**M-step**

# EM – Axis-aligned Gaussian



- We have unlabeled data  $x^1, \dots, x^m$
- We know there are  $K$  classes
- We know that the Gaussians are axis aligned

$$\Sigma_k = \begin{bmatrix} \sigma_{k1} & 0 & 0 \\ 0 & \ddots & 0 \\ 0 & 0 & \sigma_{kd} \end{bmatrix}$$

Start with an initial guess for  $\mu_1, \dots, \mu_K, \Sigma_1, \dots, \Sigma_K, \alpha_1, \dots, \alpha_K$ ,

1. Given parameters, we can easily compute probability that a point  $x^j$  belongs to class  $i$ :

$$p(y = k | x^j) \propto p(x^j | \mu_k, \Sigma_k) p(y = k)$$

**E-step**

Simply evaluate this, then normalize


2. Given the probability of each point belonging to each class, re-estimate the  $\mu_1, \dots, \mu_K, \Sigma_1, \dots, \Sigma_K, \alpha_1, \dots, \alpha_K$ ,

$$\mu_k = \frac{\sum_{j=1}^m P(y=k | x^j) x^j}{\sum_{j=1}^m P(y=k | x^j)} \quad \alpha_k = \frac{\sum_{j=1}^m P(y=k | x^j)}{m}$$

**M-step**

$$\sigma_{kl}^2 = \frac{\sum_{j=1}^m P(y=k | x^j) (x_l^j - \mu_{kl})^2}{\sum_{j=1}^m P(y=k | x^j)} \quad l\text{-th dimension}$$

# EM – General Gaussian



Start with an initial guess for  $\mu_1, \dots, \mu_K, \Sigma_1, \dots, \Sigma_K, \alpha_1, \dots, \alpha_K$ ,

1. If we know the parameters, we can easily compute probability that a point  $x^j$  belongs to class  $k$ :

$$P(y = k | x^j) \propto p(x^j | \mu_k, \Sigma_k) p(y = k)$$

**E-step**

Simply evaluate this, then normalize

2. If we know *the* probability that each point belongs to each class, we can estimate the  $\mu_1, \dots, \mu_K, \Sigma_1, \dots, \Sigma_K, \alpha_1, \dots, \alpha_K$ ,


$$\mu_k = \frac{\sum_{j=1}^m P(y=k|x^j) x^j}{\sum_{j=1}^m P(y=k|x^j)}$$

$$\alpha_k = \frac{\sum_{j=1}^m P(y = k | x^j)}{m}$$


$$\Sigma_k = \frac{\sum_{j=1}^m P(y = k | x^j) (x^j - \mu_k)(x^j - \mu_k)^T}{\sum_{j=1}^m P(y = k | x^j)}$$

**M-step**


# Gaussian Mixture Example: Start



# After first iteration




# After 2nd iteration







# After 3rd iteration




# After 4th iteration




# After 5th iteration



# After 6th iteration



# After 20th iteration



Q: Why are these two points red?


# Understanding Expectation Maximization

- EM is a general approach for learning with latent variables
  - Latent: hidden, e.g., the cluster labels
- Consider an estimation problem with a training set  $\{x_1, x_2, \dots, x_m\}$
- We wish to fit the parameter  $\theta$  of a model  $p(x, z; \theta)$  to the data without observing  $z$
- Log-likelihood function:

$$l(\theta) = \sum_{i=1}^m \log \sum_{z_i} p(x_i, z_i; \theta)$$

- Directly maximizing  $l(\theta)$  can be hard (due to the log of sum)
- EM uses a technique called **Optimization Transfer** to solve this optimization problem iteratively

# Optimization Transfer (OT)



- Given a complex function  $\Psi$  to minimize (shown as the solid line)
- OT works iteratively, minimizes a **surrogate function**  $\phi_n$  (dashed line) at each iteration:
  - $x^{n+1} = \arg \min_x \phi_n(x)$
- In any iteration  $n$ , if the surrogate function satisfy the following:
  - $\phi_n(x^n) = \Psi(x^n)$  (match at current pos)
  - $\phi_n(x) \geq \Psi(x)$  ( $\phi_n$  is the upper bound of  $\Psi$ )
- We are guaranteed to monotonically improve in each iteration
  - $\Psi(x^{n+1}) \leq \Psi(x^n)$

To apply OT for maximize, we have:

1.  $x^{n+1} = \arg \min_x \phi_n(x)$
2.  $\phi_n(x^n) = \Psi(x^n)$  (match at current pos) and  $\phi_n(x) \leq \Psi(x)$  (lower bound)

# Expectation Maximization as Optimization Transfer

- Expectation maximization uses optimization transfer to maximize the log-likelihood


$$l(\theta) = \sum_{i=1}^m \log \sum_{z_i} p(x_i, z_i; \theta)$$

- Each iteration, it finds an lower bound of the log-likelihood using **Jensen's inequality**
- It then maximizes the lower bound



# Jensen's Inequality

- Definition: a function is **convex** if the line segment between any two points on the graph of the function lies above the graph



- Jensen's inequality:**

If  $f$  is convex, and let  $x$  be a random variable, then:

$$E[f(x)] \geq f(E[x])$$

If  $f$  is concave, this is opposite and we have:  $E[f(x)] \leq f(E[x])$

# Log-likelihood and lower bound

- Objective: Log-likelihood function

Initial  $l(\theta)$

$$\begin{aligned} l(\theta) &= \sum_{i=1}^m \log \sum_{z_i} p(x_i, z_i; \theta) = \sum_{i=1}^m \log \sum_{z_i} \frac{q_i(z_i) p(x_i, z_i; \theta)}{q_i(z_i)} \\ &= \sum_{i=1}^m \log \sum_{z_i} q_i(z_i) \left[ \frac{p(x_i, z_i; \theta)}{q_i(z_i)} \right] = \sum_{i=1}^m \log E_{z_i \sim q_i(z_i)} \left[ \frac{p(x_i, z_i; \theta)}{q_i(z_i)} \right] \\ &\geq \sum_{i=1}^m E_{z_i \sim q_i(z_i)} \left[ \log \frac{p(x_i, z_i; \theta)}{q_i(z_i)} \right] \end{aligned}$$

Log is a concave function  
Using Jensen's inequality

- For any distribution  $q_i(z_i)$ , this gives a lower bound to the log-likelihood
- To be a valid surrogate for optimization transfer, we also need it to match  $l$  at current parameter  $\theta^n$ :

$$\log \sum_{z_i} q_i(z_i) \left[ \frac{p(x_i, z_i; \theta^n)}{q_i(z_i)} \right] = \sum_{z_i} q_i(z_i) \log \frac{p(x_i, z_i; \theta^n)}{q_i(z_i)}$$

# Further developing the surrogate

- We want to satisfy

$$\log \sum_{z_i} q_i(z_i) \left[ \frac{p(x_i, z_i; \theta^n)}{q_i(z_i)} \right] = \sum_{z_i} q_i(z_i) \log \frac{p(x_i, z_i; \theta^n)}{q_i(z_i)}$$

- If the circled part is constant across all possible  $z_i$  values then we will have:
  - Left side:  $\log \sum_{z_i} C q_i(z_i) = \log C \sum_{z_i} q_i(z_i) = \log C$
  - Right side:  $\sum_{z_i} q_i(z_i) \log C = \log C \sum_{z_i} q_i(z_i) = \log C$
- So we have

$$\frac{p(x_i, z_i; \theta^n)}{q_i(z_i)} = C \rightarrow q_i(z_i) = \frac{1}{C} p(x_i, z_i; \theta^n)$$

- Note that  $q_i$  must satisfy  $\sum_{z_i} q_i(z_i) = 1$

$$\frac{1}{C} \sum_{z_i} p(x_i, z_i; \theta^n) = 1$$

$$C = \sum_{z_i} p(x_i, z_i; \theta^n) = p(x_i; \theta^n)$$

$$q_i(z_i) = \frac{p(x_i, z_i; \theta^n)}{p(x_i; \theta^n)} = p(z_i | x_i; \theta^n)$$

# Expectation Maximization

Repeat until convergence {

// $\theta^n$ : current parameters in iteration  $n$

(E-step) For each data point  $i$ , compute posterior of  $z_i$ :

$$q_i(z_i) = p(z_i|x_i; \theta^n)$$

(M-step) Maximize the expected log-likelihood

$$\theta^{n+1} = \arg \max_{\theta} \sum_i \sum_{z_i} q_i(z_i) \log \frac{p(x_i, z_i; \theta)}{q_i(z_i)}$$

$$= \arg \max_{\theta} \sum_i \sum_{z_i} q_i(z_i) \log p(x_i, z_i; \theta)$$

$$n \leftarrow n + 1$$

}

# Mixture of Gaussian revisited

- Goal: given  $(x_1, \dots, x_m)$ , fitting parameters  $\alpha_1, \dots, \alpha_k; \mu_1, \dots, \mu_k; \Sigma_1, \dots, \Sigma_k$
- The cluster labels are the latent variables  $z'_i$ s
- E-step:
  - Compute  $q_i(z_i) = p(z_i|x_i; \theta^n)$  - posterior of cluster label
- M-step:
  - $\arg \max_{\theta} \sum_i \sum_{z_i} q_i(z_i) \log p(x_i, z_i|\theta)$  - maximize the expected complete loglikelihood

# Behavior of EM



- It is guaranteed to converge
  - Following the principle of Optimization Transfer
    - Each iteration it maximize a lower bound of the likelihood function
$$\theta^{n+1} = \arg \max_{\theta} \sum_i \sum_{z_i} p(z_i | x_i; \theta^n) \log p(x_i, z_i | \theta)$$
  - In practice it may converge slowly, one can stop early if the change in the log-likelihood is smaller than a threshold
- It converges to a **local optimum**
  - Multiple restart is recommended
  - Choosing the solution with the highest marginal likelihood

# Selecting k: A Model Selection Problem

- Each choice of k corresponds to a different statistical model for the data
- Model selection searches for a model ( a choice of k) that gives us the best fit of the training data
  - Penalty method
  - Cross-validation method
  - Model selection methods can also be used to make other model decisions such as choosing among different ways of constraining  $\Sigma$

# Selecting k: heuristic approaches

- For kmeans, plot the sum of squared error for different k values
  - SSE will monotonically decrease as we increase k
  - Knee points on the curve suggest possible candidates for k





# Penalty Method: Bayesian Information Criterion

- Based on Bayesian Model Selection
  - Determine the range of  $k$  values to consider  $1 \leq k \leq K_{max}$
  - Apply EM to learn a maximum likelihood fitting of the Gaussian mixture model for each possible value of  $k$
  - Choose  $k$  that maximizes BIC

$$\underbrace{2l_{\mathcal{M}}(x, \hat{\theta})}_{\text{Loglikelihood of the resulting Gaussian Mixture Model}} - \underbrace{m_{\mathcal{M}}}_{\text{\# of parameters to be estimated in } M} \log(n) \equiv \text{BIC}$$

# of data points

Loglikelihood of the resulting  
Gaussian Mixture Model

# of parameters to be estimated in  $M$

- Given two estimated models, the model with higher BIC is preferred
- Larger  $k$  increases the likelihood, but will also cause the second term to increase
- Often observed to be biased toward less complex model
- Similar method:  $\text{AIC} = 2l_m - 2m_M$ , which penalize complex model less severely

# Cross-validation Likelihood

(Smyth 1998)



- The likelihood of the training data will always increase as we increase  $k$ 
  - more clusters, more flexibility leads to better fitting of the data
- Use cross-validation
  - For each fold, learn the GMM model using the training data
  - Compute the log-likelihood of the learned model on the remaining fold as test data

# Stability Based Methods

- Stability: repeatedly produce similar clusterings on data originating from the same source. (Levine & Domany, 2001, Tibshirani and Walther, 2005)
- High level of agreement among a set of clusterings  $\Rightarrow$  the clustering model ( $k$ ) is appropriate for the data
- Evaluate multiple models, and select the model resulting in the highest level of stability.

# How to Evaluate Clustering?

- By user interpretation
  - does a document cluster seem to correspond to a specific topic?
- Internal criterion – a good clustering will produce high quality clusters:
  - high intra-cluster similarity
  - low inter-cluster similarity



- The measured quality of a clustering depends on both the object representation and the similarity measure used

# External indexes

If true class labels (*ground truth*) are known, the validity of a clustering can be verified by comparing the class labels and clustering labels.

$$\begin{array}{c|c} N & \cdot \\ \hline \cdot & n_{..} \end{array} = \begin{array}{cccc|c} n_{11} & n_{12} & \dots & n_{1l} & n_{1.} \\ n_{21} & n_{22} & \dots & n_{2l} & n_{2.} \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ n_{k1} & n_{k2} & \dots & n_{kl} & n_{k.} \\ \hline n_{.1} & n_{.2} & \dots & n_{.l} & n_{..} \end{array}$$

$n_{ij}$  = number of objects in class  $i$  and cluster  $j$

# Rand Index and Normalized Rand Index

- Given partition ( $P$ ) and ground truth ( $G$ ), measure the number of vector pairs that are:

- $a$ : in the same class both in  $P$  and  $G$ .
- $b$ : in the same class in  $P$ , but different classes in  $G$ .
- $c$ : in different classes in  $P$ , but in the same class in  $G$ .
- $d$ : in different classes both in  $P$  and  $G$ .

$$R = \frac{a + d}{a + b + c + d}$$


Rand index 1 is best. It rarely

- Adjusted rand index: corrected-for-chance version of rand index
  - Compare to the expectation of the index assuming a random partition of the same cluster sizes

$$ARI = \frac{Index - ExpectedR}{MaxIndex - ExpectedR} = \frac{\sum_{i,j} \binom{n_{ij}}{2} - \left[ \sum_i \binom{n_{i.}}{2} \sum_j \binom{n_{.j}}{2} \right] / \binom{n}{2}}{\frac{1}{2} \left[ \sum_i \binom{n_{i.}}{2} + \sum_j \binom{n_{.j}}{2} \right] - \left[ \sum_i \binom{n_{i.}}{2} \sum_j \binom{n_{.j}}{2} \right] / \binom{n}{2}}$$

# Purity and Normalized Mutual Information


- Purity



► Figure 16.1 Purity as an external evaluation criterion for cluster quality. Majority class and number of members of the majority class for the three clusters are: x, 5 (cluster 1); o, 4 (cluster 2); and  $\diamond$ , 3 (cluster 3). Purity is  $(1/17) \times (5 + 4 + 3) \approx 0.71$ .

- Normalized Mutual Information

Mutual info:



$$I(\text{Class}, \text{Clust}) = H(\text{Class}) - H(\text{Class}|\text{Clust})$$

$$NMI = \frac{2I(\text{Class}, \text{Clust})}{H(\text{Clust}) + H(\text{Class})}$$

# References for model selection

- Smyth, Padhraic. "Model selection for probabilistic clustering using cross-validated likelihood." *Statistics and Computing* 10.1 (2000): 63-72
- Erel Levine and Eytan Domany. Resampling Method for Unsupervised Estimation of Cluster Validity. *Neural Comput.* 13, 11, 2001, 2573-2593
- Tibshirani, Robert, and Guenther Walther. Cluster validation by prediction strength. *Journal of Computational and Graphical Statistics* 14.3 (2005)